


Collaborative Filtering

CS 189/289A Project T Final



Team MA
Maxwell Chen
Abinav Routhu

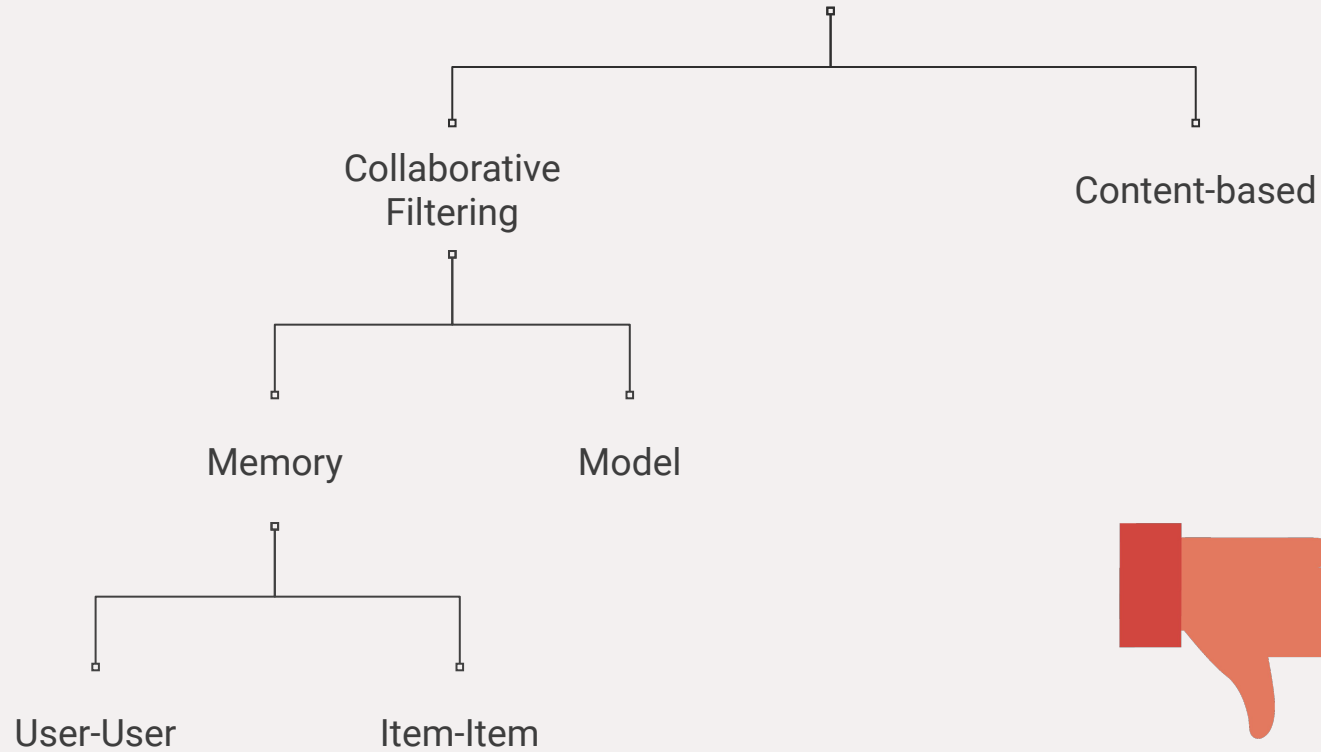
Intro

Recommendation Systems & Their Classification

Why Recommendation Systems?

- According to [McKinsey](#), 35% of Amazon.com's revenue is generated by its recommendation engine
- Netflix estimates the recommendation system saves the company around \$1 Billion annually
- Recommendations are responsible for 70% of the time people spend watching videos on YouTube

Recommendation Systems



Paradigm

Content-Based

- Require featurization
- Conceptually simple (out of box classical models)
- Phrased as regression or classification
- Difficult to exploit user-user **and** item-item relations

Collaborative Filtering

- Implicit featurization
- Novel concepts (specific to Rec. Sys.)
- Phrased as clustering or optimization problem
- Information-efficient

*Commercial deployments are overwhelmingly *hybrid* systems.

Memory Approach

The data as it is



User-Interaction Matrix

	Item 1	Item 2	Item 3	Item 4	Item 5	Item 6
User A	4	5	3	2	\emptyset	3
User B	\emptyset	\emptyset	\emptyset	5	\emptyset	\emptyset
User C	2	\emptyset	1	2	1	2
User D	\emptyset	5	\emptyset	2	\emptyset	1
User E	\emptyset	5	1	2	\emptyset	\emptyset

What would User D rate Item 3? User E rate Item 6?

- User A rates Item 2 a 5
- User B has only rated Item 4
- User C has not rated Item 2
- User C gives low ratings
- Item 4 is very popular
- Item 2 is very highly rated

- User-Interaction Matrix can be **massive** but always very **sparse** (mostly null entries)
- Sparsity not structured -- no expectation which items have been rated
- In most use cases, the # of users \gg # of items (tall matrix)
- Can decipher patterns based on similarities between users!
- Not necessary to abstract a model, can only work with data, only use what's *in memory*

0	1	0	0
3	1	3	3
1	\emptyset	3	2
\emptyset	\emptyset	\emptyset	\emptyset
4	5	3	5
\emptyset	\emptyset	4	\emptyset
2	\emptyset	1	1
\emptyset	5	\emptyset	2
\emptyset	5	1	\emptyset
3	\emptyset	\emptyset	1
\emptyset	\emptyset	\emptyset	\emptyset
\emptyset	\emptyset	\emptyset	\emptyset

User-User

1. Identify similar users to User X
2. Find highly rated items from set of similar users
3. Recommend top items not yet rated by User X

User-User

1. Identify similar users to User X

How?

Use K-NN algorithm on row vectors!

User-User

1. Identify similar users to User X
2. Find highly rated items from set of similar users

Look at column sums of submatrix.

User-User

1. Identify similar users to User X
2. Find highly rated items from set of similar users
3. Recommend top items not yet rated by User X

Check against original U-I Matrix.

User-User (In action)

User D wants
recommendations.

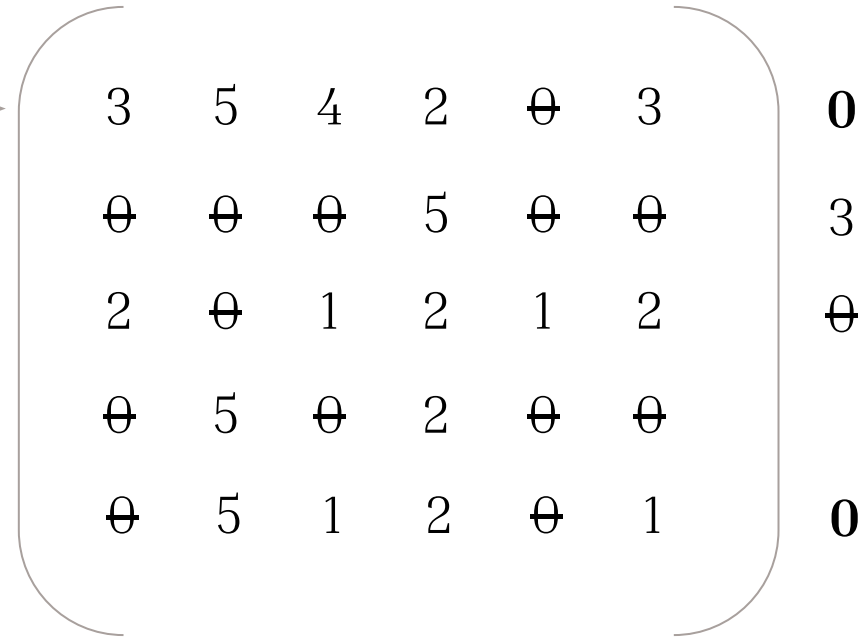


3	5	4	2	θ	3
θ	θ	θ	5	θ	θ
2	θ	1	2	1	2
θ	5	θ	2	θ	θ
θ	5	1	2	θ	1

User-User (In action)

User D is most similar
to User A and User E.

Similarity Measure is l_1 distance.
Why do we not pick User C?



3	5	4	2	\emptyset	3	0
\emptyset	\emptyset	\emptyset	5	\emptyset	\emptyset	3
2	\emptyset	1	2	1	2	\emptyset
\emptyset	5	\emptyset	2	\emptyset	\emptyset	
\emptyset	5	1	2	\emptyset	1	0

User-User (In action)

Calculate the column sums of the submatrix.

$$\begin{pmatrix} 3 & 5 & 4 & 2 & \emptyset & 3 \\ \emptyset & 5 & 1 & 2 & \emptyset & 1 \end{pmatrix}$$



3 10 5 4 0 4

User-User (In action)

User D has already
rated Item 2 and 4.

User D is
recommended Item 3
and Item 6.



3	5	4	2	\emptyset	3
\emptyset	\emptyset	\emptyset	5	\emptyset	\emptyset
2	\emptyset	1	2	1	2
\emptyset	5	\emptyset	2	\emptyset	\emptyset
\emptyset	5	1	2	\emptyset	1
3	10	5	4	0	4

User-User

Advantages

- Simple and intuitive
- Well understood
- Very personalized
- Adaptive with different similarity measures

Disadvantages

- Scales poorly because of k-nn runtime
- Unstable (very few values determine result)
- Similarity measures can have bad edge cases

Item-Item

1. Find User X's top rated items
2. Find other similar items for each item
3. Recommend most frequently found items in search

Item-Item

1. Find User X's top rated items

How?

Sort corresponding row.

Item-Item

1. Find User X's top rated items
2. Find other similar items for each item

Run k-nearest neighbors on the columns on the U-I Matrix.

Item-Item

1. Find User X's top rated items
2. Find other similar items for each item
3. Recommend most frequently found items in search

Search for repeats in list.

Item-Item (In action)

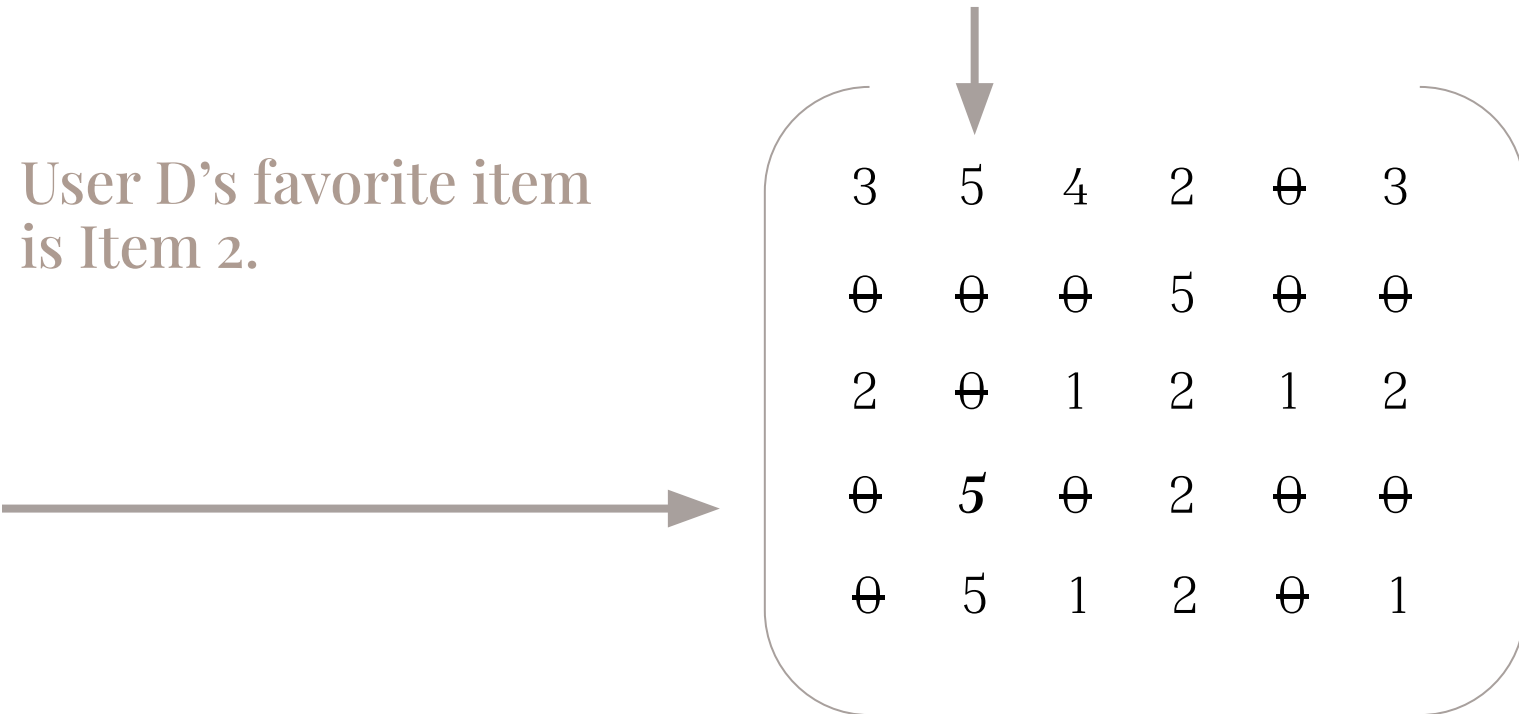
User D wants
recommendations.



3	5	4	2	θ	3
θ	θ	θ	5	θ	θ
2	θ	1	2	1	2
θ	5	θ	2	θ	θ
θ	5	1	2	θ	1

Item-Item (In action)

User D's favorite item
is Item 2.

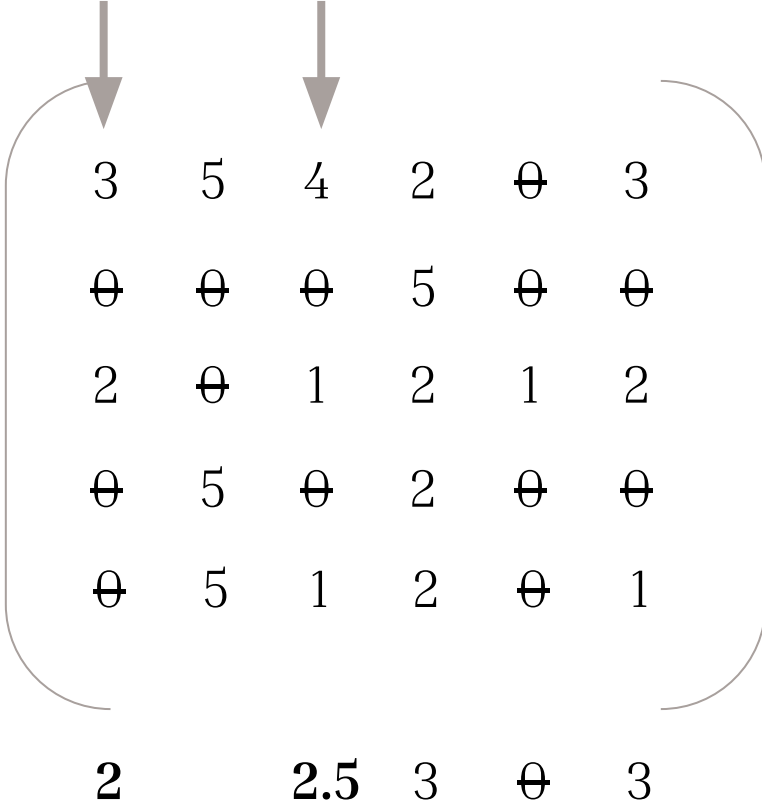


3	5	4	2	θ	3
θ	θ	θ	5	θ	θ
2	θ	1	2	1	2
θ	5	θ	2	θ	θ
θ	5	1	2	θ	1

Item-Item (In action)

Item 2 is most similar to Item 1 and Item 3.


Similarity Measure is l_1 distance.



3	5	4	2	\emptyset	3
\emptyset	\emptyset	\emptyset	5	\emptyset	\emptyset
2	\emptyset	1	2	1	2
\emptyset	5	\emptyset	2	\emptyset	\emptyset
\emptyset	5	1	2	\emptyset	1
2		2.5	3	\emptyset	3

Item-Item (In action)

User D's 2nd favorite item is Item 4.



3	5	4	2	θ	3
θ	θ	θ	5	θ	θ
2	θ	1	2	1	2
θ	5	θ	2	θ	θ
θ	5	1	2	θ	1

Item-Item (In action)

Item 4 is most similar to Item 1 and Item 6.

Similarity Measure is l_1 distance.

Why Item 6 over Item 5?

More comparisons = More reliable

3	5	4	2	\emptyset	3
\emptyset	\emptyset	\emptyset	5	\emptyset	\emptyset
2	\emptyset	1	2	1	2
\emptyset	5	\emptyset	2	\emptyset	\emptyset
\emptyset	5	1	2	\emptyset	1
0.5	3	1.3		1	1

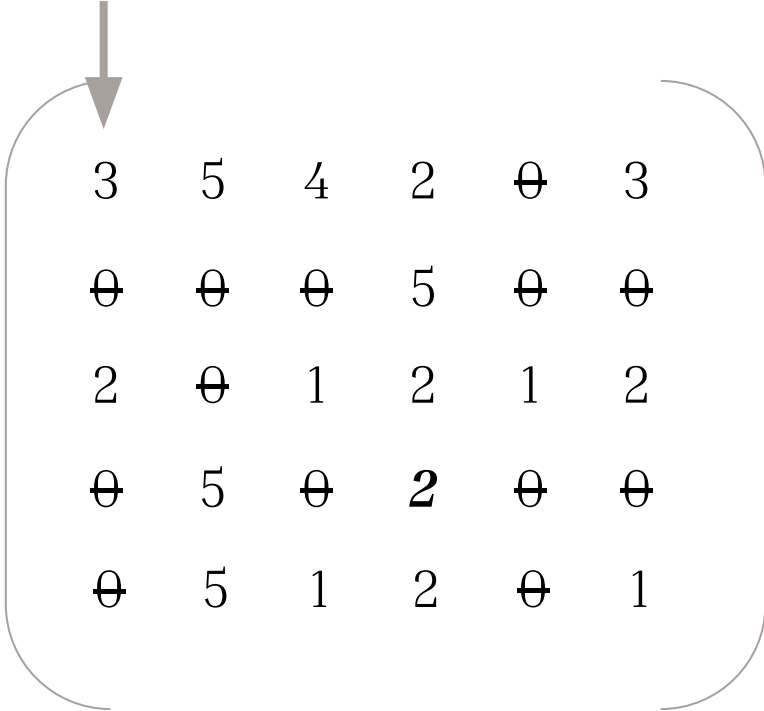
Item-Item (In action)

User D likes Item 2 and 4.

Item 2 is like Item 1 and 3.

Item 4 is like Item 1 and 6.

User D is recommended
Item **1**.



3	5	4	2	θ	3
θ	θ	θ	5	θ	θ
2	θ	1	2	1	2
θ	5	θ	2	θ	θ
θ	5	1	2	θ	1

User-User VS. Item-Item

User-User

- Simple
- Few data comparisons
→ Unstable
- Uses all of user info
→ Very personalized

Item-Item

- Also simple
- More data comparisons
→ More stable
- Uses subset of user info
→ Less specific to user

User-User VS. Item-Item

Shared Problems

- **“Cold start”**: Methods rely on known ratings.
What about new items/users with no ratings?
- **Data inefficiency**: Methods only look at subsets of User-Interaction Matrix, not efficient
- **Rich get richer**: Bias towards recommending popular items (items with many existing ratings)
- **Slow**: k-nn algorithm scales poorly
- **Similarity sensitivity**: Recommendations are *very* sensitive to choice of similarity measure

Model Approach

Latent Space Assumption & Matrix Factorization



Another Approach

What if we tried to fill in all the missing values (θ)?

4	5	3	2	θ	3
θ	θ	θ	5	θ	θ
2	θ	1	2	1	2
θ	5	θ	2	θ	θ
θ	5	1	2	θ	1

Another Approach

What if we tried to fill in all the missing values (θ)?
We need to make some sort of assumption.

4	5	3	2	θ	3
θ	θ	θ	5	θ	θ
2	θ	1	2	1	2
θ	5	θ	2	θ	θ
θ	5	1	2	θ	1

Latent Space Assumption

Assumption: The U-I matrix is *roughly* rank-1.
(From now on, we refer to the U-I matrix as **A**).

$$\mathbf{A} \approx \mathbf{U} * \mathbf{V}$$
$$\begin{pmatrix} 4 & 5 & 3 & 2 & \mathbf{2} & 3 \\ \mathbf{1} & \mathbf{3} & \mathbf{2} & 5 & \mathbf{5} & \mathbf{4} \\ 2 & \mathbf{1} & 1 & 2 & 1 & 2 \\ \mathbf{4} & 5 & \mathbf{3} & 2 & \mathbf{1} & \mathbf{2} \\ \mathbf{2} & 5 & 1 & 2 & \mathbf{4} & 1 \end{pmatrix} \approx \begin{pmatrix} -7.9 & 1.9 \\ -8.0 & -3.8 \\ -3.5 & -0.3 \\ -7.2 & 2.6 \\ -6.7 & -0.3 \end{pmatrix} \begin{pmatrix} -0.4 & -0.6 & -0.3 & -0.4 & -0.4 & -0.4 \\ 0.5 & 0.4 & 0.2 & -0.4 & -0.6 & -0.2 \end{pmatrix}$$

Latent Space Assumption

$$\begin{array}{c} \xrightarrow{\text{A}} \\ \xrightarrow{\text{B}} \\ \xrightarrow{\text{C}} \\ \xrightarrow{\text{D}} \\ \xrightarrow{\text{E}} \end{array} \begin{pmatrix} -7.9 & 1.9 \\ -8.0 & -3.8 \\ -3.5 & -0.3 \\ -7.2 & 2.6 \\ -6.7 & -0.3 \end{pmatrix}$$

$$\begin{array}{cccccc} 1 & 2 & 3 & 4 & 5 & 6 \\ \downarrow & \downarrow & \downarrow & \downarrow & \downarrow & \downarrow \\ \begin{pmatrix} -0.4 & -0.6 & -0.3 & -0.4 & -0.4 & -0.4 \\ 0.5 & 0.4 & 0.2 & -0.4 & -0.6 & -0.2 \end{pmatrix} \end{array}$$

\mathbf{U} - user embeddings

\mathbf{V} - item embeddings

Latent Space Assumption

Side Note

How to generate \mathbf{U} and \mathbf{V} ?

(Eckhart-Young Theorem)

If \mathbf{A} was known fully (no missing ratings),
compact SVD of \mathbf{A} gives best rank- l approximation.

Not very useful since goal is to predict missing ratings in first place...

Instead, we can try to \mathbf{U} , \mathbf{V} via gradient descent in so-called “SVD” algorithm (that, confusingly, does not use the SVD)

Matrix Factorization

Measure accuracy of \mathbf{U} , \mathbf{V} by known entries in \mathbf{A} !

$$\begin{pmatrix} 4 & 5 & 3 & 2 & & 3 \\ & & & 5 & & \\ 2 & & 1 & 2 & 1 & 2 \\ & 5 & & 2 & & \\ & 5 & 1 & 2 & & 1 \end{pmatrix} \approx \mathbf{U} * \mathbf{V}$$

$$\begin{pmatrix} \mathbf{4.1} & \mathbf{5.5} & \mathbf{2.8} & \mathbf{2.4} & \mathbf{2.0} & \mathbf{2.8} \\ 1.3 & 3.3 & 1.6 & \mathbf{4.7} & 5.5 & 3.9 \\ \mathbf{1.3} & 2.0 & \mathbf{1.0} & \mathbf{1.5} & \mathbf{1.6} & \mathbf{1.5} \\ 4.2 & \mathbf{5.4} & 2.7 & \mathbf{1.8} & 1.3 & 2.4 \\ \mathbf{2.5} & \mathbf{3.9} & \mathbf{2.0} & \mathbf{2.8} & 2.9 & \mathbf{2.7} \end{pmatrix} = \begin{pmatrix} -7.9 & 1.9 \\ -8.0 & -3.8 \\ -3.5 & -0.3 \\ -7.2 & 2.6 \\ -6.7 & -0.3 \end{pmatrix} \begin{pmatrix} -0.4 & -0.6 & -0.3 & -0.4 & -0.4 & -0.4 \\ 0.5 & 0.4 & 0.2 & -0.4 & -0.6 & -0.2 \end{pmatrix}$$

Matrix Factorization

Measure accuracy of U, V by known entries in A !

$$\min \left\| \begin{pmatrix} 4 & 5 & 3 & 2 & & 3 \\ & & & 5 & & \\ 2 & & 1 & 2 & 1 & 2 \\ & 5 & & 2 & & \\ & 5 & 1 & 2 & & 1 \end{pmatrix} - \begin{pmatrix} 4.1 & 5.5 & 2.8 & 2.4 & 2.0 & 2.8 \\ 1.3 & 3.3 & 1.6 & 4.7 & 5.5 & 3.9 \\ 1.3 & 2.0 & 1.0 & 1.5 & 1.6 & 1.5 \\ 4.2 & 5.4 & 2.7 & 1.8 & 1.3 & 2.4 \\ 2.5 & 3.9 & 2.0 & 2.8 & 2.9 & 2.7 \end{pmatrix} \right\|$$



$$\min \quad \frac{1}{2} \sum (a_{ij} - u_i^T v_j)^2 \quad \text{where } a_{ij} \in K$$

Matrix Factorization

Measure accuracy of \mathbf{U} , \mathbf{V} by known entries in \mathbf{A} !

$$\min \quad \frac{1}{2} \sum (\mathbf{a}_{ij} - \mathbf{u}_i^T \mathbf{v}_j)^2 \quad \text{where } \mathbf{a}_{ij} \in K$$

\mathbf{a}_{ij} : entry of \mathbf{A} at row i and col j

\mathbf{u}_i^T : row i of \mathbf{U} \mathbf{v}_j : col j of \mathbf{V}

K : set of all known (rated) entries of \mathbf{A}

Matrix Factorization

$$\min \frac{1}{2} \sum (\mathbf{a}_{ij} - \mathbf{u}_i^T \mathbf{v}_j)^2 \quad \text{where } \mathbf{a}_{ij} \in K$$

Main Weakness:

Has many parameters (nml) and is highly prone to overfit
(high variance, low bias)

Solution:

Have lots of data, introduce regularization, and use a “baseline” model

Matrix Factorization (Improved)

$$\min \quad \frac{1}{2} \sum (\mathbf{a}_{ij} - (\mu + \mathbf{b}_j^v + \mathbf{b}_i^u + \mathbf{u}_i^T \mathbf{v}_j))^2 + \lambda (\mathbf{b}_j^v + \mathbf{b}_i^u + \|\mathbf{u}_i\|^2 + \|\mathbf{v}_j\|^2)$$

where $\mathbf{a}_{ij} \in K$

\mathbf{a}_{ij} : entry of \mathbf{A} at row i and col j

\mathbf{u}_i^T : row i of \mathbf{U} \mathbf{v}_j : col j of \mathbf{V}

K : set of all known (rated) entries of \mathbf{A}

μ : global average of known entries of \mathbf{A}

\mathbf{b}_j^v : offset for item j \mathbf{b}_i^u : rating offset for user i

λ : ridge coefficient

Matrix Factorization (Improved)

$$\min \quad \frac{1}{2} \sum (\mathbf{a}_{ij} - (\mu + \mathbf{b}_j^v + \mathbf{b}_i^u + \mathbf{u}_i^T \mathbf{v}_j))^2 + \lambda (\mathbf{b}_j^v + \mathbf{b}_i^u + \|\mathbf{u}_i\|^2 + \|\mathbf{v}_j\|^2)$$

where $\mathbf{a}_{ij} \in K$

What did we change?

- 1) Our estimate of \mathbf{a}_{ij} used to simply be $\mathbf{u}_i^T \mathbf{v}_j$: the embedding product.
Now, we estimate it as $\mu + \mathbf{b}_j^v + \mathbf{b}_i^u + \mathbf{u}_i^T \mathbf{v}_j$: the global rating mean + item j 's offset + user i 's offset + embedding product.
- 2) We now penalize our estimate to prevent it from growing too big.

Matrix Factorization (Improved)

Rationale behind change #1

Suppose, we wished to estimate user i 's rating of item j .

1. Starting guess would just be an average rating overall (μ)
2. What if item j is a bad product, generally reviewed poorly? ($b_j^v < 0$)
3. But suppose user i is very generous with her ratings? ($b_i^u > 0$)
4. Result? Abstract the biases for user i , item j , and the rating system overall. Want to isolate what the embedding measures: user i 's specific preference for item j

Matrix Factorization (Improved)

$$\min \quad \frac{1}{2} \sum (\mathbf{a}_{ij} - (\mu + \mathbf{b}_j^v + \mathbf{b}_i^u + \mathbf{u}_i^T \mathbf{v}_j))^2 + \lambda (\mathbf{b}_j^v + \mathbf{b}_i^u + \|\mathbf{u}_i\|^2 + \|\mathbf{v}_j\|^2)$$

where $\mathbf{a}_{ij} \in K$

Congrats! This is the celebrated “SVD” algorithm.

- Developed in *Netflix Prize* contest, machine learning challenge to build best collab filtering algorithm with \$1M prize (Author won 3rd)
- Later refined into SVD++ algorithm by adding 1 more term
- Outdated, current state of the art uses deep learning

The End